

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import requests
import tweepy
import json
import re
import os
import sys
```

Data Gathering, Part 1: Twitter Archive

```
In [2]: df1 = pd.read_csv('twitter-archive-enhanced.csv')
```

Twitter Archive Visual Assessment

```
In [3]: df1.shape
```

```
Out[3]: (2356, 17)
```

```
In [4]: df1.columns
```

```
Out[4]: Index(['tweet_id', 'in_reply_to_status_id', 'in_reply_to_user_id', 'timestamp',
              'source', 'text', 'retweeted_status_id', 'retweeted_status_user_id',
              'retweeted_status_timestamp', 'expanded_urls', 'rating_numerator',
              'rating_denominator', 'name', 'doggo', 'floofer', 'pupper', 'pupper'],
              dtype='object')
```

In [6]: df1.head(8)

Out[6]:

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	timestamp	
0	8.924210e+17	NaN	NaN	2017-08-01 16:23:56 +0000	href="http://twitter.com/downl
1	8.921770e+17	NaN	NaN	2017-08-01 00:17:27 +0000	href="http://twitter.com/downl
2	8.918150e+17	NaN	NaN	2017-07-31 00:18:03 +0000	href="http://twitter.com/downl
3	8.916900e+17	NaN	NaN	2017-07-30 15:58:51 +0000	href="http://twitter.com/downl
4	8.913280e+17	NaN	NaN	2017-07-29 16:00:24 +0000	href="http://twitter.com/downl
5	8.910880e+17	NaN	NaN	2017-07-29 00:08:17 +0000	href="http://twitter.com/downl
6	8.909720e+17	NaN	NaN	2017-07-28 16:27:12 +0000	href="http://twitter.com/downl
7	8.907290e+17	NaN	NaN	2017-07-28 00:22:40 +0000	href="http://twitter.com/downl

In [7]: `df1.tail()`

Out[7]:

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	timestamp	
2351	6.660490e+17	NaN	NaN	2015-11-16 00:24:50 +0000	href="http://twitter.com/do
2352	6.660440e+17	NaN	NaN	2015-11-16 00:04:52 +0000	href="http://twitter.com/do
2353	6.660330e+17	NaN	NaN	2015-11-15 23:21:54 +0000	href="http://twitter.com/do
2354	6.660290e+17	NaN	NaN	2015-11-15 23:05:30 +0000	href="http://twitter.com/do
2355	6.660210e+17	NaN	NaN	2015-11-15 22:32:08 +0000	href="http://twitter.com/do

Twitter Archive Programmatic Assessment

In [8]: `# Let's check how many unique tweet ID's are included in this data set.
df1['tweet_id'].nunique()`

Out[8]: 2349

In [9]: df1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null float64
in_reply_to_status_id   78 non-null float64
in_reply_to_user_id     78 non-null float64
timestamp               2356 non-null object
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id     181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls           2297 non-null object
rating_numerator        2356 non-null int64
rating_denominator      2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
dtypes: float64(5), int64(2), object(10)
memory usage: 313.0+ KB
```

In [12]: *# Let's make sure that the 'doggo' and 'floofer' columns only have 'None'*
 print(df1.doggo.unique())
 print(df1.floofer.unique())

```
['None' 'doggo']
['None' 'floofer']
```

```
In [13]: print(df1.doggo.value_counts())
print(df1.floofer.value_counts())
print(df1.pupper.value_counts())
print(df1.puppo.value_counts())
```

```
None      2259
doggo      97
Name: doggo, dtype: int64
None      2346
floofer    10
Name: floofer, dtype: int64
None      2099
pupper     257
Name: pupper, dtype: int64
None      2326
puppo      30
Name: puppo, dtype: int64
```

```
In [14]: df1.duplicated().sum()
```

```
Out[14]: 0
```

```
In [15]: # The 'text' column appears to be the only column with no duplicates
df1['text'].duplicated().sum()
```

```
Out[15]: 0
```

Twitter Archive Cleaning and Tidiness Issues Found:

Data Quality:

- tweet_id needs to be converted to a string
- retweeted columns need to be converted to strings
- timestamp columns need to be converted to a datetime format
- names such as 'a' and 'such' should be changed to 'None'
- in the 'text' column the url should be separated from the actual tweet
- some of the rating numerators are way outside of the range that we would expect so this will need to be corrected
- some of the rating denominators are numbers other than 10 so we will need to set them to 10
- some of the 'text' columns indicate a dog name but that name is missing from the 'name' column

Tidiness:

- we may want to melt the 'doggo', 'pupper', etc. columns into a single column

Data Gathering, Part 2: Image Predictions archive

```
In [16]: # Import of image predictions file.
predicted_breeds_url = 'https://d17h27t6h515a5.cloudfront.net/topher/201
response = requests.get(predicted_breeds_url)
with open('image_predictions.tsv', 'wb') as f:
    f.write(response.content)
```

```
In [17]: df2 = pd.read_csv('image_predictions.tsv', sep='\t')
```

The following items are explanations of the image predictions as provided by Udacity:

- tweet_id is the last part of the tweet URL after "status/" → https://twitter.com/dog_rates/status/889531135344209921 (https://twitter.com/dog_rates/status/889531135344209921)
- p1 is the algorithm's #1 prediction for the image in the tweet → golden retriever
- p1_conf is how confident the algorithm is in its #1 prediction → 95%
- p1_dog is whether or not the #1 prediction is a breed of dog → TRUE
- p2 is the algorithm's second most likely prediction → Labrador retriever
- p2_conf is how confident the algorithm is in its #2 prediction → 1%
- p2_dog is whether or not the #2 prediction is a breed of dog → TRUE

Image Predictions Visual Assessment

In [18]: df2.shape

Out[18]: (2075, 12)

In [19]: df2.head(8)

Out[19]:

	tweet_id	jpg_url	img_num	
0	666020888022790149	https://pbs.twimg.com/media/CT4udn0WwAA0aMy.jpg	1	Welsh_spring
1	666029285002620928	https://pbs.twimg.com/media/CT42GRgUYAA5iDo.jpg	1	
2	666033412701032449	https://pbs.twimg.com/media/CT4521TWwAEvMyu.jpg	1	German
3	666044226329800704	https://pbs.twimg.com/media/CT5Dr8HUEAA-IEu.jpg	1	Rhodesian_
4	666049248165822465	https://pbs.twimg.com/media/CT5lQmsXIAAKY4A.jpg	1	miniature
5	666050758794694657	https://pbs.twimg.com/media/CT5Jof1WUAEuVxN.jpg	1	Bernese_mou
6	666051853826850816	https://pbs.twimg.com/media/CT5KoJ1WoAAJash.jpg	1	
7	666055525042405380	https://pbs.twimg.com/media/CT5N9tpXIAAifs1.jpg	1	

Image Predictions Programmatic Assessment

```
In [20]: print(df2.p1_dog.value_counts())
print(df2.p2_dog.value_counts())
print(df2.p3_dog.value_counts())
```

```
True      1532
False      543
Name: p1_dog, dtype: int64
True      1553
False      522
Name: p2_dog, dtype: int64
True       1499
False       576
Name: p3_dog, dtype: int64
```

```
In [21]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null int64
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
p2            2075 non-null object
p2_conf       2075 non-null float64
p2_dog        2075 non-null bool
p3            2075 non-null object
p3_conf       2075 non-null float64
p3_dog        2075 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB
```



```
In [22]: # I would like to look at what particular dog breeds appear multiple times
print(df2.p1.value_counts()[0:5])
print(df2.p2.value_counts()[0:5])
print(df2.p3.value_counts()[0:5])
```

```
golden_retriever      150
Labrador_retriever    100
Pembroke              89
Chihuahua             83
pug                   57
Name: p1, dtype: int64
Labrador_retriever    104
golden_retriever      92
Cardigan              73
Chihuahua             44
Pomeranian           42
Name: p2, dtype: int64
Labrador_retriever     79
Chihuahua             58
golden_retriever       48
Eskimo_dog            38
kelpie                35
Name: p3, dtype: int64
```

```
In [23]: # Note that all tweet_id's in this DataFrame are unique.
df2.tweet_id.nunique()
```

Out[23]: 2075

Image Predictions Cleaning and Tidiness Issues Found:

Data Quality:

- the image prediction number 'img_num' should be converted to a string

Tidiness:

- we will need to merge this DataFrame with df1

Data Gathering, Part 3: Retweets and Likes via Twitter API

```
In [14]: import tweepy
consumer_key = '#####'
consumer_secret = '#####'
access_token = '#####&&&&&'
access_secret = '#####'

auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
auth.set_access_token(access_token, access_secret)

api = tweepy.API(auth)
```

The following StackOverflow link was helpful for obtaining likes and retweets:

<https://stackoverflow.com/questions/45761253/how-do-i-get-the-number-of-likes-on-a-tweet-via-tweepy> (<https://stackoverflow.com/questions/45761253/how-do-i-get-the-number-of-likes-on-a-tweet-via-tweepy>)

We will create short lists with each tweet's favorites and retweets and then use a dictionary format to create a third DataFrame named *df3*.

```
In [15]: # samples:
tweet0 = api.get_status(df2['tweet_id'][0])
tweet1 = api.get_status(df2['tweet_id'][1])
print(tweet0.retweet_count, tweet0.favorite_count, '\n')
print(tweet1.retweet_count, tweet1.favorite_count, '\n')
```

498 2531

46 126

The following sections of code pull the counts of retweets and favorites for the Tweet ID's in our Image Predictions file. I found it necessary to write multiple iterations of the same code in order to pull the data little by little and avoid a rate limit error from the Twitter API.

```
In [16]: retweet_list1 = []
for i in df2['tweet_id'][0:300]:
    try:
        tweet_info = api.get_status(i)
        retweet_list1.append(tweet_info.retweet_count)
    except:
        retweet_list1.append(0)
```

```
In [17]: retweet_list2 = []
         for i in df2['tweet_id'][300:600]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list2.append(tweet_info.retweet_count)
             except:
                 retweet_list2.append(0)
```

```
In [18]: retweet_list3 = []
         for i in df2['tweet_id'][600:900]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list3.append(tweet_info.retweet_count)
             except:
                 retweet_list3.append(0)
```

```
In [19]: retweet_list4 = []
         for i in df2['tweet_id'][900:1200]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list4.append(tweet_info.retweet_count)
             except:
                 retweet_list4.append(0)
```

```
In [20]: retweet_list5 = []
         for i in df2['tweet_id'][1200:1500]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list5.append(tweet_info.retweet_count)
             except:
                 retweet_list5.append(0)
```

```
In [21]: retweet_list6 = []
         for i in df2['tweet_id'][1500:1800]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list6.append(tweet_info.retweet_count)
             except:
                 retweet_list6.append(0)
```

```
In [37]: retweet_list7 = []
for i in df2['tweet_id'][1800:]:
    try:
        tweet_info = api.get_status(i)
        retweet_list7.append(tweet_info.retweet_count)
    except:
        retweet_list7.append(0)
```

```
In [64]: retweet_complete = retweet_list1 + retweet_list2 + retweet_list3 + retweet_list4 + retweet_list5 + retweet_list6 + retweet_list7
```

```
In [39]: favorite_list1 = []
for i in df2['tweet_id'][0:300]:
    try:
        tweet_info = api.get_status(i)
        favorite_list1.append(tweet_info.favorite_count)
    except:
        favorite_list1.append(0)
```

```
In [47]: favorite_list2 = []
for i in df2['tweet_id'][300:600]:
    try:
        tweet_info = api.get_status(i)
        favorite_list2.append(tweet_info.favorite_count)
    except:
        favorite_list2.append(0)
```

```
In [48]: favorite_list3 = []
for i in df2['tweet_id'][600:900]:
    try:
        tweet_info = api.get_status(i)
        favorite_list3.append(tweet_info.favorite_count)
    except:
        favorite_list3.append(0)
```

```
In [49]: favorite_list4 = []
for i in df2['tweet_id'][900:1200]:
    try:
        tweet_info = api.get_status(i)
        favorite_list4.append(tweet_info.favorite_count)
    except:
        favorite_list4.append(0)
```

```
In [50]: favorite_list5 = []
        for i in df2['tweet_id'][1200:1500]:
            try:
                tweet_info = api.get_status(i)
                favorite_list5.append(tweet_info.favorite_count)
            except:
                favorite_list5.append(0)
```

```
In [62]: favorite_list6 = []
        for i in df2['tweet_id'][1500:1800]:
            try:
                tweet_info = api.get_status(i)
                favorite_list6.append(tweet_info.favorite_count)
            except:
                favorite_list6.append(0)
```

```
In [63]: favorite_list7 = []
        for i in df2['tweet_id'][1800:]:
            try:
                tweet_info = api.get_status(i)
                favorite_list7.append(tweet_info.favorite_count)
            except:
                favorite_list7.append(0)
```

```
In [65]: favorites_complete = favorite_list1 + favorite_list2 + favorite_list3 +
        favorite_list6 + favorite_list7
```

```
In [69]: df3 = pd.DataFrame({'tweet_id': df2['tweet_id'], 'retweet_count': retweet_count})
```

```
In [72]: df2_tweet_ids = df2['tweet_id'].tolist()
```

```
In [73]: likes_dictionary = {'tweet_id': df2_tweet_ids, 'retweet_count': retweet_count}
```

I found the following article helpful in working with json files and dictionaries:

<https://stackoverflow.com/questions/42825102/how-to-save-python-dictionary-into-json-files>
<https://stackoverflow.com/questions/42825102/how-to-save-python-dictionary-into-json-files>

```
In [74]: # Let's store this information in a JSON text file.
        jsonarray = json.dumps(likes_dictionary)
```

```
In [75]: with open('tweet_json.txt', "w") as fp:
        json.dump(jsonarray, fp)
```

I stored the retweet and favorite data in an Excel sheet in addition to the JSON .txt file. Also, for ease of use there is a read_excel file below so that the data can be brought in quickly and easily of the Jupyter Notebook is closed and restarted.

```
In [25]: df3.to_excel('retweet_output.xlsx', index=False)
```

```
In [26]: df3 = pd.read_excel('retweet_output.xlsx')
```

Twitter API Data Visual and Programmatic Assessment

```
In [27]: df3.shape
```

```
Out[27]: (2075, 3)
```

```
In [28]: df3.columns
```

```
Out[28]: Index(['tweet_id', 'retweet_count', 'favorite_count'], dtype='object')
```

```
In [30]: df3.tail()
```

```
Out[30]:
```

	tweet_id	retweet_count	favorite_count
2070	891327558926688256	9120	39425
2071	891689557279858688	8421	41254
2072	891815181378084864	4051	24519
2073	892177421306343424	6115	32558
2074	892420643555336192	8278	37900

```
In [31]: # Let's check the averages.
df3['retweet_count'].mean(), df3['favorite_count'].mean()
```

```
Out[31]: (2777.757108433735, 8342.715662650602)
```

Data Cleaning

```
In [32]: df1_clean = df1.copy()
df2_clean = df2.copy()
df3_clean = df3.copy()
```

Cleaning step 1: Convert tweet_id in all three DataFrames to a string

```
In [33]: df1_clean['tweet_id'] = df1_clean['tweet_id'].astype(str)
df2_clean['tweet_id'] = df2_clean['tweet_id'].astype(str)
df3_clean['tweet_id'] = df3_clean['tweet_id'].astype(str)
```

Cleaning step 2: Convert 'retweeted' columns in df1 to a strings

```
In [35]: df1_clean['retweeted_status_id'] = df1_clean['retweeted_status_id'].astype(str)
df1_clean['retweeted_status_user_id'] = df1_clean['retweeted_status_user_id'].astype(str)
```

Cleaning step 3: Convert timestamps from a string to a time format

```
In [36]: df1_clean['timestamp'] = pd.to_datetime(df1_clean['timestamp'])
df1_clean['retweeted_status_timestamp'] = pd.to_datetime(df1_clean['retweeted_status_timestamp'])
```

```
In [37]: # Let's test the first three cleaning items above.
df1_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null object
in_reply_to_status_id    78 non-null float64
in_reply_to_user_id      78 non-null float64
timestamp               2356 non-null datetime64[ns]
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id      2356 non-null object
retweeted_status_user_id 2356 non-null object
retweeted_status_timestamp 181 non-null datetime64[ns]
expanded_urls            2297 non-null object
rating_numerator         2356 non-null int64
rating_denominator       2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
dtypes: datetime64[ns](2), float64(2), int64(2), object(11)
memory usage: 313.0+ KB
```

Cleaning step 4: Convert image number from an integer to a string

```
In [25]: df2_clean['img_num'] = df2_clean['img_num'].astype(str)
```

```
In [38]: # Verification:
df2_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null object
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
p2            2075 non-null object
p2_conf       2075 non-null float64
p2_dog        2075 non-null bool
p3            2075 non-null object
p3_conf       2075 non-null float64
p3_dog        2075 non-null bool
dtypes: bool(3), float64(3), int64(1), object(5)
memory usage: 152.1+ KB
```

Cleaning step 5: Update some of the names in the 'name' column

```
In [41]: df1_clean['name'].replace(['a', 'an', 'such', 'the', 'quite'], 'None', inplace=True)
```

```
In [50]: # Verification of the code above. The terms replaced should be zero and
#        than 1:

df1_clean.query("name == 'a'")['name'].count(), df1_clean.query("name == 'an'")['name'].count(),
df1_clean.query("name == 'the'")['name'].count(), df1_clean.query("name == 'quite'")['name'].count(),
df1_clean.query("name == 'Charlie'")['name'].count()
```

```
Out[50]: (0, 0, 0, 0, 12)
```

Cleaning step 6: Split the weblink section of the 'text' column into its own separate column.
This will be helpful if we search for terms or do any other specific work on the 'text' column.

```
In [51]: df1_clean['tweet_link'] = df1_clean['text'].str.extract(r'(https://t.co/')
```

```
In [52]: df4 = df1_clean.text.str.partition(' https://')
```



```
In [53]: df4.shape
```

```
Out[53]: (2356, 3)
```

```
In [54]: df4.rename(columns={0: 'text_section'}, inplace=True)
df4.drop([1, 2], axis=1, inplace=True)
```

```
In [55]: # Let's add this 'text_section' column to df1.
df1_new = pd.concat([df1_clean, df4], axis=1)
df1_new.drop('text', axis=1, inplace=True)
```

```
In [57]: # Verification:
df1_new.columns
```

```
Out[57]: Index(['tweet_id', 'in_reply_to_status_id', 'in_reply_to_user_id', 'timestamp',
               'source', 'retweeted_status_id', 'retweeted_status_user_id',
               'retweeted_status_timestamp', 'expanded_urls', 'rating_numerator',
               'rating_denominator', 'name', 'doggo', 'floofer', 'pupper', 'puppo',
               'tweet_link', 'text_section'],
              dtype='object')
```

Cleaning Step 7: Set the rating numerators to a certain minimum and maximum. After reviewing a number of entries in the 'text' column it appears that the minimum can remain at zero, but the maximum should be set to 15.

```
In [59]: print('current min:', min(df1_new['rating_numerator']))
print('current max:', max(df1_new['rating_numerator']))
```

```
current min: 0
current max: 1776
```

```
In [60]: rating_num = df1_new['rating_numerator'].tolist()
```

```
In [61]: rating_numerator_updated = [min(x, 15) for x in rating_num]
```

```
In [62]: # We will verify this step and the next one below.
df1_new['rating_numerator_clean'] = rating_numerator_updated
df1_new.drop(['rating_numerator'], axis=1, inplace=True)
```

Cleaning Step 8: Set the rating denominators to 10.

```
In [63]: df1_new.rating_denominator.value_counts()
```

```
Out[63]: 10      2333
         11         3
         50         3
         80         2
         20         2
          2         1
         16         1
         40         1
         70         1
         15         1
         90         1
        110         1
        120         1
        130         1
        150         1
        170         1
          7         1
          0         1
        Name: rating_denominator, dtype: int64
```

```
In [64]: rating_denom = df1_new['rating_denominator'].tolist()
```

```
In [65]: df1_new['rating_denominator_clean'] = df1_new['rating_denominator'].where
```

```
In [66]: df1_new.rating_denominator_clean.value_counts()
```

```
Out[66]: 10      2356
        Name: rating_denominator_clean, dtype: int64
```

```
In [67]: df1_new.drop('rating_denominator', axis=1, inplace=True)
```

```
In [69]: # Verification of steps 8 and 9:
df1_new.rating_denominator_clean.value_counts(), df1_new.rating_numerator_clean.value_counts()
```

```
Out[69]: (10      2356
          Name: rating_denominator_clean, dtype: int64, 12      558
          11      464
          10      461
          13      351
           9      158
           8      102
           7       55
          14       54
           5       37
           6       32
          15       28
           3       19
           4       17
           1        9
           2        9
           0         2
          Name: rating_numerator_clean, dtype: int64)
```

Cleaning Step 9: Obtain correct tweet_id's. The tweet_id's in the 'tweet_id' column of the twitter-archive-enhanced are incomplete. In order to join or merge our data frames on the tweet_id column we will need to have consistency in this column.

```
In [70]: df1_new['tweet_id'] = df1_new['expanded_urls'].str.extract(r'(\d{18})')
```

Cleaning Step 10: Let's see if we can make some updates to the 'name' column since there are a number of cases where the dog's name can be found in the 'text' column but not the name column. At the very least I would like to see if the dog is named or not since that will be useful in our analysis later.

```
In [71]: name_in_text = []

for text in range(df1_new.shape[0]):
    try:
        # what are we searching for?
        word1 = re.compile(r'name(.*)')
        # where are we searching?
        mol = word1.search(df1_new.text_section[text])
        name1 = mol.group()
        name_in_text.append(name1)
    except:
        name_in_text.append('')
```

```
In [72]: df1_new['name2'] = name_in_text
```

```
In [73]: has_name = []

for i in range(df1_new.shape[0]):
    if (df1_new['name'][i] == 'None') & (df1_new['name2'][i] == ''):
        has_name.append('No')
    else:
        has_name.append('Yes')
```

```
In [74]: df1_new['dog_has_name'] = has_name
```

```
In [75]: # Verification:
df1_new['dog_has_name'].value_counts()
```

```
Out[75]: Yes      1571
        No       785
        Name: dog_has_name, dtype: int64
```

```
In [76]: df1_new.drop('name2', axis=1, inplace=True)
```

```
In [77]: # Final cleaning verification:
df1_new.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 19 columns):
tweet_id                2190 non-null object
in_reply_to_status_id   78 non-null float64
in_reply_to_user_id     78 non-null float64
timestamp               2356 non-null datetime64[ns]
source                  2356 non-null object
retweeted_status_id     2356 non-null object
retweeted_status_user_id 2356 non-null object
retweeted_status_timestamp 181 non-null datetime64[ns]
expanded_urls           2297 non-null object
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
tweet_link              2231 non-null object
text_section            2356 non-null object
rating_numerator_clean  2356 non-null int64
rating_denominator_clean 2356 non-null int64
dog_has_name            2356 non-null object
dtypes: datetime64[ns](2), float64(2), int64(2), object(13)
memory usage: 349.8+ KB
```

Tidiness step 1:

Combine the 'doggo', 'floofer', 'pupper' and 'puppo' columns into a single column.

```
In [78]: df1_new['dog_labels'] = df1_new['doggo'] + df1_new['floofer'] + df1_new[
```

```
In [79]: df1_new['dog_labels'].value_counts()
```

```
Out[79]: NoneNoneNoneNone      1976
NoneNonepupperNone      245
doggoNoneNoneNone       83
NoneNoneNonepuppo       29
doggoNonepupperNone      12
NoneflooferNoneNone       9
doggoflooferNoneNone      1
doggoNoneNonepuppo        1
Name: dog_labels, dtype: int64
```

```
In [80]: df1_new['dog_labels'] = df1_new.dog_labels.str.split('None')
```

```
In [81]: df1_new['dog_labels'] = df1_new.dog_labels.apply(', '.join)
```

The following Stackoverflow article was helpful:

<https://stackoverflow.com/questions/37347725/convert-a-panda-df-list-into-a-string>
(<https://stackoverflow.com/questions/37347725/convert-a-panda-df-list-into-a-string>)

```
In [82]: df1_new['dog_labels'] = df1_new.dog_labels.str.strip(',')
```

```
In [83]: df1_new['dog_labels'] = np.where(df1_new['dog_labels'] == 'doggo,,puppo'
df1_new['dog_labels'] = np.where(df1_new['dog_labels'] == 'doggo,pupper'
df1_new['dog_labels'] = np.where(df1_new['dog_labels'] == 'doggofloofer'
```

```
In [84]: # verification step
df1_new.dog_labels.value_counts()
```

```
Out[84]:
```

	1976
pupper	245
doggo	83
puppo	29
doggo and pupper	12
floofer	9
doggo and puppo	1
doggo and floofer	1
Name: dog_labels, dtype: int64	

```
In [85]: df1_new.drop(['doggo', 'floofer', 'pupper', 'puppo'], axis=1, inplace=True)
```

Tidiness step 2:

Combine 3 DataFrames into 1

```
In [86]: print(df1_new.shape)
print(df2_clean.shape)
print(df3_clean.shape)
```

```
(2356, 16)
(2075, 12)
(2075, 3)
```

```
In [87]: df2_clean.sort_values('tweet_id', inplace=True)
```

```
In [88]: df3_clean.rename(columns={'tweet_id': 'tweet_id2'}, inplace=True)
```

```
In [89]: df3_clean.sort_values('tweet_id2', inplace=True)
```

```
In [90]: df5 = pd.concat([df2_clean, df3_clean], axis=1)
```

```
In [91]: # Check to make sure that the first fifteen digits of the tweet_id are e
id_check = []
for a in range(df2_clean.shape[0]):
    if df5.tweet_id[a][0:15] == df5.tweet_id2[a][0:15]:
        id_check.append('OK')
    else:
        id_check.append('no_match')
```

```
In [92]: df5['tweet_id_check'] = id_check
```

```
In [93]: df5.tweet_id_check.value_counts()
```

```
Out[93]: OK      2075
Name: tweet_id_check, dtype: int64
```

```
In [94]: df5.drop('tweet_id_check', axis=1, inplace=True)
```

```
In [95]: df1_new.sort_values('tweet_id', ascending=True, inplace=True)
```

```
In [96]: # Let's do a visual check to see that the first few tweet id's match up
print(df1_new.tweet_id.head())
print(df5.tweet_id.head())
```

```
2355    666020888022790149
2354    666029285002620928
2353    666033412701032449
2352    666044226329800704
2351    666049248165822465
Name: tweet_id, dtype: object
0      666020888022790149
1      666029285002620928
2      666033412701032449
3      666044226329800704
4      666049248165822465
Name: tweet_id, dtype: object
```

```
In [97]: df6 = pd.merge(df1_new, df5, on='tweet_id', how='inner', indicator=True)
```

```
In [98]: df6.shape
```

```
Out[98]: (2138, 31)
```

```
In [99]: print(df6.tweet_id.duplicated().sum())
```

```
144
```

Note: df6 will be our master dataset for analysis and visualization

```
In [100]: df6.to_csv('twitter_archive_master.csv', index=False)
```

For each of the insights below, a thorough explanation is provided separately in the 'wrapgle_report' document.

Insights, Part 1: Relationship between Ratings and Number of Likes (the favorite_count).

```
In [101]: # Let's make a scatterplot to view the relationship between a dog's rating and the number of likes it has.
# First we need to remove the 144 duplicate tweets from our dataset.
df6A = df6.drop_duplicates('tweet_id', keep='last')
print(df6A.shape)
```

```
(1994, 31)
```

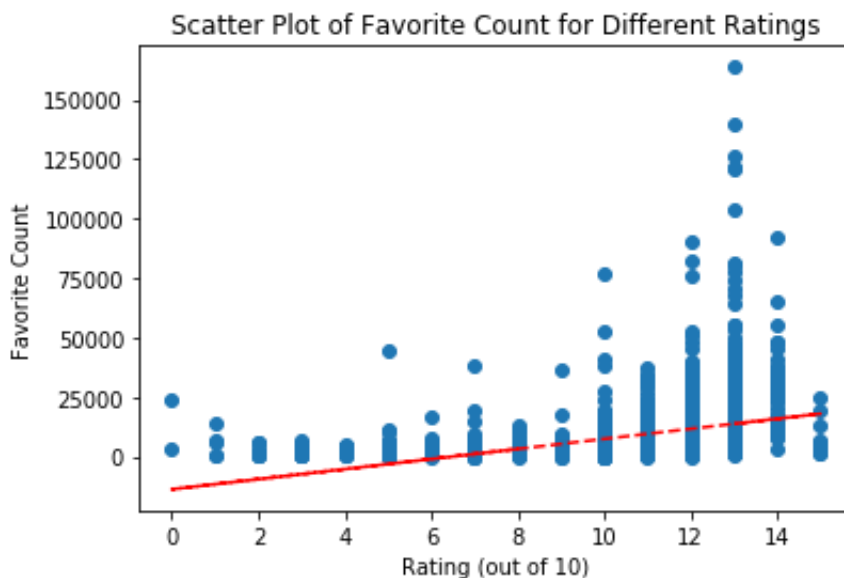
The following Stack Overflow article was helpful in setting up a trendline:

<https://stackoverflow.com/questions/41635448/how-can-i-draw-scatter-trend-line-on-matplotlib-python-pandas?noredirect=1> (<https://stackoverflow.com/questions/41635448/how-can-i-draw-scatter-trend-line-on-matplotlib-python-pandas?noredirect=1>)


```
In [102]: x = df6A.rating_numerator_clean
y = df6A.favorite_count
plt.scatter(x, y)
plt.title('Scatter Plot of Favorite Count for Different Ratings')
plt.xlabel('Rating (out of 10)')
plt.ylabel('Favorite Count')

z = np.polyfit(x, y, 1)
p = np.poly1d(z)
plt.plot(x, p(x), "r--")

plt.show()
```



Insights, Part 2: Named Dogs vs Unnamed Dogs

My impression is that it is easier for viewers to connect with a dog with a name, so we will test to see if they have more likes on average.

As with Part 1 of our data analysis we should use df6A so that duplicates are removed.

Item 10 in the data cleaning section identifies whether or not the dog has a name.

```
In [103]: yes_name = df6A.query("dog_has_name == 'Yes'")['rating_numerator_clean']
no_name = df6A.query("dog_has_name == 'No'")['rating_numerator_clean'].m

# next we need the averages so we can take the sum using a groupby and t
```

```
In [104]: print(yes_name)
          print(no_name)
```

```
10.710021321961621
10.270868824531517
```

Insights, Part 3: A Look at the Use of the Dog Label Terms

For this last analysis section I would like to get a general comparison at how often the terms 'doggo', 'floofer', 'puppo' and 'pupper' are used in the 'text' column and how often they are used in the labels columns that was included with the data.

```
In [105]: df6A_rows = df6A.shape[0]
          doggo_in_text = []
          floofer_in_text = []
          pupper_in_text = []
          puppo_in_text = []
          empty_list = []
          label_list = []

          for text in range(df6A_rows):
              try:
                  # what are we searching for?
                  word1 = re.compile(r'doggo\s')
                  # where are we searching?
                  finding1 = word1.findall(df6A.text_section[text])
                  label_list.append(finding1)
                  word2 = re.compile(r'floofer\s')
                  finding2 = word2.findall(df6A.text_section[text])
                  label_list.append(finding2)
                  word3 = re.compile(r'pupper\s')
                  finding3 = word3.findall(df6A.text_section[text])
                  label_list.append(finding3)
                  word4 = re.compile(r'puppo\s')
                  finding4 = word4.findall(df6A.text_section[text])
                  label_list.append(finding4)
              except:
                  empty_list.append('')
```

```
In [106]: print('doggo:', label_list.count(['doggo ']))
print('floofer:', label_list.count(['floofer ']))
print('pupper:', label_list.count(['pupper ']))
print('puppo:', label_list.count(['puppo ']))
```

```
doggo: 40
floofer: 0
pupper: 75
puppo: 5
```

```
In [107]: # Let's compare the figures above to the number of times they appears in
df6A['dog_labels'].value_counts()
```

```
Out[107]:
pupper      203
doggo       63
puppo       22
doggo and pupper    9
floofer       7
doggo and puppo     1
doggo and floofer   1
Name: dog_labels, dtype: int64
```

Insights, Part 4: Dog Breeds and Image Predictions

The neural network used to predict the dog breeds often predicted Chihuahua, Labrador retriever, and golden retriever. Let's compare the accuracy for each of those three breeds on the first two attempts.

```
In [108]: print(df6A.query("p1 == 'Chihuahua'")['p1_dog'].value_counts())
print(df6A.query("p1 == 'Labrador_retriever'")['p1_dog'].value_counts())
print(df6A.query("p1 == 'golden_retriever'")['p1_dog'].value_counts())
```

```
True      79
Name: p1_dog, dtype: int64
True      95
Name: p1_dog, dtype: int64
True     139
Name: p1_dog, dtype: int64
```

```
In [109]: print(df6A.query("p2 == 'Chihuahua'")['p2_dog'].value_counts())
print(df6A.query("p2 == 'Labrador_retriever'")['p2_dog'].value_counts())
print(df6A.query("p2 == 'golden_retriever'")['p2_dog'].value_counts())
```

```
True      43
Name: p2_dog, dtype: int64
True      96
Name: p2_dog, dtype: int64
True      82
Name: p2_dog, dtype: int64
```