

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import requests
import tweepy
import json
import re
import os
import sys
```

## Data Gathering, Part 1: Twitter archive

```
In [2]: df1 = pd.read_csv('twitter-archive-enhanced.csv')
```

```
In [3]: df1.shape
```

```
Out[3]: (2356, 17)
```

```
In [4]: # Let's check how many unique tweet ID's are included in this data set
df1['tweet_id'].nunique()
```

```
Out[4]: 2349
```

```
In [5]: df1.columns
```

```
Out[5]: Index(['tweet_id', 'in_reply_to_status_id', 'in_reply_to_user_id', 'timestamp',
              'source', 'text', 'retweeted_status_id', 'retweeted_status_user_id',
              'retweeted_status_timestamp', 'expanded_urls', 'rating_numerator',
              'rating_denominator', 'name', 'doggo', 'floofer', 'pupper', 'puppo'],
              dtype='object')
```

In [6]: `df1.head()`

Out[6]:

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	timestamp	
0	8.924210e+17	NaN	NaN	2017-08-01 16:23:56 +0000	<a href="http://twit r...
1	8.921770e+17	NaN	NaN	2017-08-01 00:17:27 +0000	<a href="http://twit r...
2	8.918150e+17	NaN	NaN	2017-07-31 00:18:03 +0000	<a href="http://twit r...
3	8.916900e+17	NaN	NaN	2017-07-30 15:58:51 +0000	<a href="http://twit r...
4	8.913280e+17	NaN	NaN	2017-07-29 16:00:24 +0000	<a href="http://twit r...

In [7]: df1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null float64
in_reply_to_status_id   78 non-null float64
in_reply_to_user_id     78 non-null float64
timestamp               2356 non-null object
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id     181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls           2297 non-null object
rating_numerator        2356 non-null int64
rating_denominator      2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
dtypes: float64(5), int64(2), object(10)
memory usage: 313.0+ KB
```

In [8]: print(df1.doggo.unique())  
print(df1.floofer.unique())

```
['None' 'doggo']
['None' 'floofer']
```

In [9]: print(df1.doggo.value\_counts())  
print(df1.floofer.value\_counts())  
print(df1.pupper.value\_counts())  
print(df1.puppo.value\_counts())

```
None      2259
doggo      97
Name: doggo, dtype: int64
None      2346
floofer    10
Name: floofer, dtype: int64
None      2099
pupper     257
Name: pupper, dtype: int64
None      2326
puppo      30
Name: puppo, dtype: int64
```

```
In [10]: df1.duplicated().sum()
```

```
Out[10]: 0
```

```
In [11]: # The 'text' column appears to be the only column with no duplicates.
df1['text'].duplicated().sum()
```

```
Out[11]: 0
```

## Data Gathering, Part 2: Image Predictions archive

```
In [12]: # Import of image predictions file.
predicted_breeds_url = 'https://d17h27t6h515a5.cloudfront.net/topher/2017/August/599fd2ad_image-predictions/image-predictions.tsv'
response = requests.get(predicted_breeds_url)
with open('image_predictions.tsv', 'wb') as f:
    f.write(response.content)
```

```
In [13]: df2 = pd.read_csv('image_predictions.tsv', sep='\t')
```

The following items are explanations of the image predictions as provided by Udacity:

- tweet\_id is the last part of the tweet URL after "status/" → [https://twitter.com/dog\\_rates/status/889531135344209921](https://twitter.com/dog_rates/status/889531135344209921) ([https://twitter.com/dog\\_rates/status/889531135344209921](https://twitter.com/dog_rates/status/889531135344209921))
- p1 is the algorithm's #1 prediction for the image in the tweet → golden retriever
- p1\_conf is how confident the algorithm is in its #1 prediction → 95%
- p1\_dog is whether or not the #1 prediction is a breed of dog → TRUE
- p2 is the algorithm's second most likely prediction → Labrador retriever
- p2\_conf is how confident the algorithm is in its #2 prediction → 1%
- p2\_dog is whether or not the #2 prediction is a breed of dog → TRUE

```
In [14]: df2.shape
```

```
Out[14]: (2075, 12)
```

```
In [15]: print(df2.p1_dog.value_counts())
print(df2.p2_dog.value_counts())
print(df2.p3_dog.value_counts())
```

```
True      1532
False      543
Name: p1_dog, dtype: int64
True      1553
False      522
Name: p2_dog, dtype: int64
True      1499
False      576
Name: p3_dog, dtype: int64
```

```
In [16]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null int64
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
p2            2075 non-null object
p2_conf       2075 non-null float64
p2_dog        2075 non-null bool
p3            2075 non-null object
p3_conf       2075 non-null float64
p3_dog        2075 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB
```

```
In [17]: # I would like to look at what particular dog breeds appear multiple times. This may be useful in our analysis.
print(df2.p1.value_counts()[0:5])
print(df2.p2.value_counts()[0:5])
print(df2.p3.value_counts()[0:5])
```

```
golden_retriever      150
Labrador_retriever    100
Pembroke              89
Chihuahua             83
pug                   57
Name: p1, dtype: int64
Labrador_retriever    104
golden_retriever      92
Cardigan              73
Chihuahua             44
Pomeranian           42
Name: p2, dtype: int64
Labrador_retriever     79
Chihuahua             58
golden_retriever      48
Eskimo_dog            38
kelpie                35
Name: p3, dtype: int64
```

```
In [18]: # Note that all tweet_id's in this DataFrame are unique.
df2.tweet_id.nunique()
```

Out[18]: 2075

```
In [19]: df2.head()
```

Out[19]:

	tweet_id	jpg_url	img_num
0	666020888022790149	https://pbs.twimg.com/media/CT4udn0WwAA0aMy.jpg	1
1	666029285002620928	https://pbs.twimg.com/media/CT42GRgUYAA5iDo.jpg	1
2	666033412701032449	https://pbs.twimg.com/media/CT4521TWwAEvMyu.jpg	1
3	666044226329800704	https://pbs.twimg.com/media/CT5Dr8HUEAA-IEu.jpg	1
4	666049248165822465	https://pbs.twimg.com/media/CT5IQmsXIAAKY4A.jpg	1

## Data Gathering, Part 3: Retweets and Likes via Twitter API

```
In [19]: import tweepy
consumer_key = 'yIAEcJfJ7uPrQKnI0zD8e6nMP'
consumer_secret = '9oVI1iwvHuVWhaejSIAurNzN5MXPr4L7nphuKjmlB2eEOngNwp'
access_token = '3220892162-QGX3QS3oRxQufzvDupEcz7EvfA2LnoQCgDchZDj'
access_secret = 'DrZ05DvXJ9GpEgLvfeKloraCqBC80zFXghMpUiTfnoazG'

auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
auth.set_access_token(access_token, access_secret)

api = tweepy.API(auth)
```

The following StackOverflow link was helpful for obtaining likes and retweets:

<https://stackoverflow.com/questions/45761253/how-do-i-get-the-number-of-likes-on-a-tweet-via-tweepy>  
(<https://stackoverflow.com/questions/45761253/how-do-i-get-the-number-of-likes-on-a-tweet-via-tweepy>)

```
In [10]: # samples:
tweet0 = api.get_status(df2['tweet_id'][0])
tweet1 = api.get_status(df2['tweet_id'][1])
print(tweet0.retweet_count, tweet0.favorite_count, '\n')
print(tweet1.retweet_count, tweet1.favorite_count, '\n')

499 2532

47 125
```

The following sections of code pull the counts of retweets and favorites for the Tweet ID's in our Image Predictions file. I found it necessary to write multiple iterations of the same code in order to pull the data little by little and avoid a rate limit error from the Twitter API.

```
In [11]: retweet_list1 = []
for i in df2['tweet_id'][0:300]:
    try:
        tweet_info = api.get_status(i)
        retweet_list1.append(tweet_info.retweet_count)
    except:
        retweet_list1.append(0)
```

```
In [13]: retweet_list2 = []
         for i in df2['tweet_id'][300:600]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list2.append(tweet_info.retweet_count)
             except:
                 retweet_list2.append(0)
```

```
In [22]: retweet_list3 = []
         for i in df2['tweet_id'][600:900]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list3.append(tweet_info.retweet_count)
             except:
                 retweet_list3.append(0)
```

```
In [28]: retweet_list4 = []
         for i in df2['tweet_id'][900:1200]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list4.append(tweet_info.retweet_count)
             except:
                 retweet_list4.append(0)
```

```
In [33]: retweet_list5 = []
         for i in df2['tweet_id'][1200:1500]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list5.append(tweet_info.retweet_count)
             except:
                 retweet_list5.append(0)
```

```
In [39]: retweet_list6 = []
         for i in df2['tweet_id'][1500:1800]:
             try:
                 tweet_info = api.get_status(i)
                 retweet_list6.append(tweet_info.retweet_count)
             except:
                 retweet_list6.append(0)
```



```
In [50]: retweet_list7 = []
for i in df2['tweet_id'][1800:]:
    try:
        tweet_info = api.get_status(i)
        retweet_list7.append(tweet_info.retweet_count)
    except:
        retweet_list7.append(0)
```

```
In [61]: retweet_complete = retweet_list1 + retweet_list2 + retweet_list3 + retweet_list4 + retweet_list5 + retweet_list6 + \
        retweet_list7
        # pd.to_excel('')
```

```
In [59]: # this list is created as a separate section of code in order to avoid exceeding the rate limit
        favorite_list1 = []
        for i in df2['tweet_id'][0:300]:
            try:
                tweet_info = api.get_status(i)
                favorite_list1.append(tweet_info.favorite_count)
            except:
                favorite_list1.append(0)
```

```
In [64]: favorite_list2 = []
        for i in df2['tweet_id'][300:600]:
            try:
                tweet_info = api.get_status(i)
                favorite_list2.append(tweet_info.favorite_count)
            except:
                favorite_list2.append(0)
```

```
In [68]: favorite_list3 = []
        for i in df2['tweet_id'][600:900]:
            try:
                tweet_info = api.get_status(i)
                favorite_list3.append(tweet_info.favorite_count)
            except:
                favorite_list3.append(0)
```

```
In [70]: favorite_list4 = []
        for i in df2['tweet_id'][900:1200]:
            try:
                tweet_info = api.get_status(i)
                favorite_list4.append(tweet_info.favorite_count)
            except:
                favorite_list4.append(0)
```

```
In [103]: favorite_list5 = []
          for i in df2['tweet_id'][1200:1500]:
              try:
                  tweet_info = api.get_status(i)
                  favorite_list5.append(tweet_info.favorite_count)
              except:
                  favorite_list5.append(0)
```

```
In [95]: favorite_list6 = []
          for i in df2['tweet_id'][1500:1800]:
              try:
                  tweet_info = api.get_status(i)
                  favorite_list6.append(tweet_info.favorite_count)
              except:
                  favorite_list6.append(0)
```

```
In [109]: favorite_list7 = []
           for i in df2['tweet_id'][1800:]:
               try:
                   tweet_info = api.get_status(i)
                   favorite_list7.append(tweet_info.favorite_count)
               except:
                   favorite_list7.append(0)
```

```
In [111]: favorites_complete = favorite_list1 + favorite_list2 + favorite_list3
          + favorite_list4 + favorite_list5 + \
            favorite_list6 + favorite_list7
```

```
In [134]: df3 = pd.DataFrame({'tweet_id': df2['tweet_id'], 'retweet_count': retw
eet_complete, 'favorite_count': favorites_complete})
```

```
In [135]: df3.shape
```

```
Out[135]: (2075, 3)
```

```
In [65]: df3.tail()
```

Out[65]:

	tweet_id	retweet_count	favorite_count
<b>2070</b>	891327558926688256	9130	39484
<b>2071</b>	891689557279858688	8428	41290
<b>2072</b>	891815181378084864	4054	24537
<b>2073</b>	892177421306343424	6122	32588
<b>2074</b>	892420643555336192	8289	37949

Note that the retweet and favorite data will be stored in an Excel sheet so that the code above does not need to be rerun each time this Jupyter Notebook is closed and reopened.

```
In [137]: df3.to_excel('retweet_output.xlsx', index=False)
```

```
In [20]: df3 = pd.read_excel('retweet_output.xlsx')
```

## Data Cleaning

```
In [21]: df1_clean = df1.copy()
df2_clean = df2.copy()
df3_clean = df3.copy()
```

Cleaning step 1: Convert tweet\_id in all three DataFrames to a string

```
In [22]: df1_clean['tweet_id'] = df1_clean['tweet_id'].astype(str)
df2_clean['tweet_id'] = df2_clean['tweet_id'].astype(str)
df3_clean['tweet_id'] = df3_clean['tweet_id'].astype(str)
```

Cleaning step 2: Convert 'retweeted' columns in df1 to a strings

```
In [23]: df1_clean['retweeted_status_id'] = df1_clean['retweeted_status_id'].as
type(str)
df1_clean['retweeted_status_user_id'] = df1_clean['retweeted_status_us
er_id'].astype(str)
```

Cleaning step 3: Convert timestamps from a string to a time format

```
In [24]: df1_clean['timestamp'] = pd.to_datetime(df1_clean['timestamp'])
df1_clean['retweeted_status_timestamp'] = pd.to_datetime(df1_clean['retweeted_status_timestamp'])
```

Cleaning step 4: Convert image number from an integer to a string

```
In [25]: df2_clean['img_num'] = df2_clean['img_num'].astype(str)
```

Cleaning step 5: Update some of the names in the 'name' column

```
In [26]: df1_clean['name'].replace(['a', 'an', 'such', 'the', 'quite'], 'None',
inplace=True)
```

Cleaning step 6: Split the weblink section of the 'text' column into its own separate column. This will be helpful if we search for terms or do any other specific work on the 'text' column.

```
In [27]: df1_clean['tweet_link'] = df1_clean['text'].str.extract(r'(https://t.co/\w*)')
```

```
In [28]: df4 = df1_clean.text.str.partition(' https://')
```

```
In [29]: df4.shape
```

```
Out[29]: (2356, 3)
```

```
In [30]: df4.rename(columns={0: 'text_section'}, inplace=True)
df4.drop([1, 2], axis=1, inplace=True)
```

```
In [31]: # Let's add this 'text_section' column to df1.
df1_new = pd.concat([df1_clean, df4], axis=1)
df1_new.drop('text', axis=1, inplace=True)
```

Cleaning Step 7: Set the rating numerators to a certain minimum and maximum. After reviewing a number of entries in the 'text' column it appears that the minimum can remain at zero, but the maximum should be set to 15.

```
In [32]: print(min(df1_new['rating_numerator']))
print(max(df1_new['rating_numerator']))

0
1776
```

```
In [33]: rating_num = df1_new['rating_numerator'].tolist()
```

```
In [34]: rating_numerator_updated = [min(x, 15) for x in rating_num]
```

```
In [35]: df1_new['rating_numerator_clean'] = rating_numerator_updated
df1_new.drop(['rating_numerator'], axis=1, inplace=True)
```

Cleaning Step 8: Set the rating denominators to 10.

```
In [36]: df1_new.rating_denominator.value_counts()
```

```
Out[36]: 10      2333
11         3
50         3
80         2
20         2
2          1
16         1
40         1
70         1
15         1
90         1
110        1
120        1
130        1
150        1
170        1
7          1
0          1
Name: rating_denominator, dtype: int64
```

```
In [37]: rating_denom = df1_new['rating_denominator'].tolist()
```

```
In [38]: df1_new['rating_denominator_clean'] = df1_new['rating_denominator'].where(df1_new['rating_denominator'] == 10, 10)
```

```
In [39]: df1_new.rating_denominator_clean.value_counts()
```

```
Out[39]: 10      2356
         Name: rating_denominator_clean, dtype: int64
```

```
In [40]: df1_new.drop('rating_denominator', axis=1, inplace=True)
```

Cleaning Step 9: Obtain correct tweet\_id's. The tweet\_id's in the 'tweet\_id' column of the twitter-archive-enhanced are incomplete. In order to join or merge our data frames on the tweet\_id column we will need to have consistency in this column.

```
In [41]: df1_new['tweet_id'] = df1_new['expanded_urls'].str.extract(r'(\d{18})')
         )
```

Cleaning Step 10: Let's see if we can make some updates to the 'name' column since there are a number of cases where the dog's name can be found in the 'text' column but not the name column. At the very least I would like to see if the dog is named or not since that will be useful in our analysis later.

```
In [42]: name_in_text = []

         for text in range(df1_new.shape[0]):
             try:
                 # what are we searching for?
                 word1 = re.compile(r'name(.*)')
                 # where are we searching?
                 mol = word1.search(df1_new.text_section[text])
                 name1 = mol.group()
                 name_in_text.append(name1)
             except:
                 name_in_text.append('')
```

```
In [43]: df1_new['name2'] = name_in_text
```

```
In [44]: has_name = []

         for i in range(df1_new.shape[0]):
             if (df1_new['name'][i] == 'None') & (df1_new['name2'][i] == ''):
                 has_name.append('No')
             else:
                 has_name.append('Yes')
```

```
In [45]: df1_new['dog_has_name'] = has_name
```

```
In [46]: df1_new['dog_has_name'].value_counts()
```

```
Out[46]: Yes      1571
        No       785
        Name: dog_has_name, dtype: int64
```

```
In [47]: df1_new.drop('name2', axis=1, inplace=True)
```

### Tidiness step 1:

Combine the 'doggo', 'floofer', 'pupper' and 'puppo' columns into a single column.

```
In [48]: df1_new['dog_labels'] = df1_new['doggo'] + df1_new['floofer'] + df1_new['pupper'] + df1_new['puppo']
```

```
In [50]: df1_new['dog_labels'].value_counts()
```

```
Out[50]: NoneNoneNoneNone      1976
        NoneNonepupperNone      245
        doggoNoneNoneNone       83
        NoneNoneNonepuppo       29
        doggoNonepupperNone      12
        NoneflooferNoneNone       9
        doggoNoneNonepuppo        1
        doggoflooferNoneNone       1
        Name: dog_labels, dtype: int64
```

```
In [51]: df1_new['dog_labels'] = df1_new.dog_labels.str.split('None')
```

```
In [52]: df1_new['dog_labels'] = df1_new.dog_labels.apply(', '.join)
```

The following Stackoverflow article was helpful: <https://stackoverflow.com/questions/37347725/converting-a-panda-df-list-into-a-string> (<https://stackoverflow.com/questions/37347725/converting-a-panda-df-list-into-a-string>)

```
In [53]: df1_new['dog_labels'] = df1_new.dog_labels.str.strip(',')
```

```
In [54]: df1_new['dog_labels'] = np.where(df1_new['dog_labels'] == 'doggo,,pupp
o', 'doggo and puppo', df1_new['dog_labels'])
df1_new['dog_labels'] = np.where(df1_new['dog_labels'] == 'doggo,puppe
r', 'doggo and pupper', df1_new['dog_labels'])
df1_new['dog_labels'] = np.where(df1_new['dog_labels'] == 'doggofloofe
r', 'doggo and floofer', df1_new['dog_labels'])
```

```
In [55]: # verification step
df1_new.dog_labels.value_counts()
```

```
Out[55]:
```

	1976
pupper	245
doggo	83
puppo	29
doggo and pupper	12
floofer	9
doggo and puppo	1
doggo and floofer	1

Name: dog\_labels, dtype: int64

```
In [56]: df1_new.drop(['doggo', 'floofer', 'pupper', 'puppo'], axis=1, inplace=
True)
```

## Tidiness step 2:

Combine 3 DataFrames into 1

```
In [57]: print(df1_new.shape)
print(df2_clean.shape)
print(df3_clean.shape)

(2356, 16)
(2075, 12)
(2075, 3)
```

```
In [58]: df2_clean.sort_values('tweet_id', inplace=True)
```

```
In [59]: df3_clean.rename(columns={'tweet_id': 'tweet_id2'}, inplace=True)
```

```
In [60]: df3_clean.sort_values('tweet_id2', inplace=True)
```

```
In [61]: df5 = pd.concat([df2_clean, df3_clean], axis=1)
```



```
In [62]: # Check to make sure that the first fifteen digits of the tweet_id are
equal in all rows.
id_check = []
for a in range(df2_clean.shape[0]):
    if df5.tweet_id[a][0:15] == df5.tweet_id2[a][0:15]:
        id_check.append('OK')
    else:
        id_check.append('no_match')
```

```
In [63]: df5['tweet_id_check'] = id_check
```

```
In [64]: df5.tweet_id_check.value_counts()
```

```
Out[64]: OK      2075
Name: tweet_id_check, dtype: int64
```

```
In [65]: df5.drop('tweet_id_check', axis=1, inplace=True)
```

```
In [66]: df1_new.sort_values('tweet_id', ascending=True, inplace=True)
```

```
In [67]: # Let's do a visual check to see that the first few tweet id's match u
p before merging the DataFrames.
print(df1_new.tweet_id.head())
print(df5.tweet_id.head())
```

```
2355      666020888022790149
2354      666029285002620928
2353      666033412701032449
2352      666044226329800704
2351      666049248165822465
Name: tweet_id, dtype: object
0      666020888022790149
1      666029285002620928
2      666033412701032449
3      666044226329800704
4      666049248165822465
Name: tweet_id, dtype: object
```

```
In [68]: df6 = pd.merge(df1_new, df5, on='tweet_id', how='inner', indicator=True
e)
```

```
In [69]: df6.shape
```

```
Out[69]: (2138, 31)
```

```
In [70]: print(df6.tweet_id.duplicated().sum())
144
```

**Note: df6 will be our master dataset for analysis and visualization**

```
In [71]: df6.to_csv('master_dataset.csv', index=False)
```

## Data Analysis, Part 1: Relationship between Ratings and Number of Likes (the favorite\_count).

```
In [72]: # Let's make a scatterplot to view the relationship between a dog's rating and the number of likes it received.
# First we need to remove the 144 duplicate tweets from our dataset.
df6A = df6.drop_duplicates('tweet_id', keep='last')
print(df6A.shape)

(1994, 31)
```

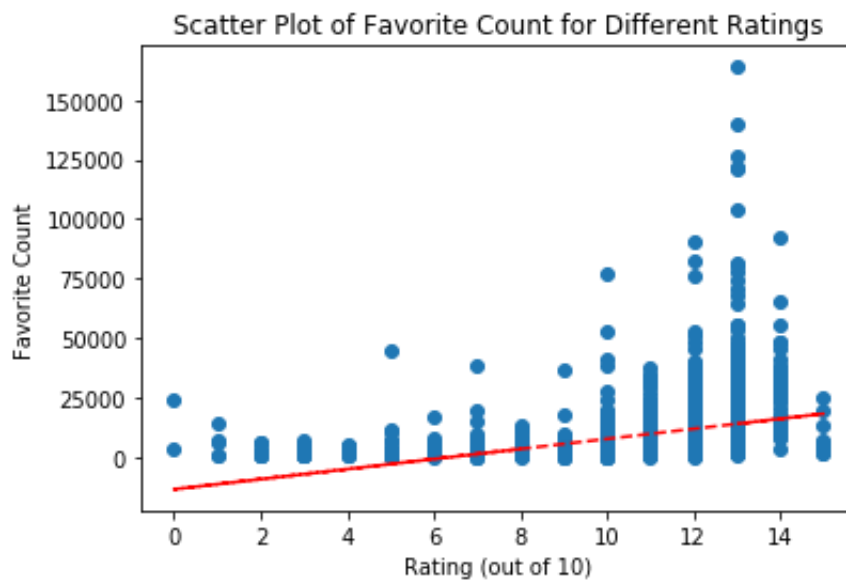
The following Stack Overflow article was helpful in setting up a trendline:

<https://stackoverflow.com/questions/41635448/how-can-i-draw-scatter-trend-line-on-matplot-python-pandas?noredirect=1> (<https://stackoverflow.com/questions/41635448/how-can-i-draw-scatter-trend-line-on-matplot-python-pandas?noredirect=1>)

```
In [73]: x = df6A.rating_numerator_clean
y = df6A.favorite_count
plt.scatter(x, y)
plt.title('Scatter Plot of Favorite Count for Different Ratings')
plt.xlabel('Rating (out of 10)')
plt.ylabel('Favorite Count')

z = np.polyfit(x, y, 1)
p = np.poly1d(z)
plt.plot(x, p(x), "r--")

plt.show()
```



The scatterplot above indicates that, in general, tweets with a higher dog rating received more likes.

## Data Analysis, Part 2: Named Dogs vs Unnamed Dogs

My impression is that it is easier for viewers to connect with a dog with a name, so we will test to see if they have more likes on average.

As with Part 1 of our data analysis we should use df6A so that duplicates are removed.

Item 10 in the data cleaning section identifies whether or not the dog has a name.

```
In [74]: yes_name = df6A.query("dog_has_name == 'Yes'")['rating_numerator_clean']
         no_name = df6A.query("dog_has_name == 'No'")['rating_numerator_clean']

         # next we need the averages so we can take the sum using a groupby and
         # then dividing by the variables above
```

```
In [75]: print(yes_name)
         print(no_name)

10.710021321961621
10.270868824531517
```

These two averages are quite close so it is reasonable to conclude that the presence of a name in the text of the tweet does NOT have a meaningful impact on how many favorites a tweet receives.

## Data Analysis, Part 3: A Look at the Use of the Dog Label Terms

For this last analysis section I would like to get a general comparison at how often the terms 'doggo', 'floofer', 'puppo' and 'pupper' are used in the 'text' column and how often they are used in the labels columns that was included with the data.

```
In [76]: df6A_rows = df6A.shape[0]
doggo_in_text = []
floofer_in_text = []
pupper_in_text = []
puppo_in_text = []
empty_list = []
label_list = []

for text in range(df6A_rows):
    try:
        # what are we searching for?
        word1 = re.compile(r'doggo\s')
        # where are we searching?
        finding1 = word1.findall(df6A.text_section[text])
        label_list.append(finding1)
        word2 = re.compile(r'floofer\s')
        finding2 = word2.findall(df6A.text_section[text])
        label_list.append(finding2)
        word3 = re.compile(r'pupper\s')
        finding3 = word3.findall(df6A.text_section[text])
        label_list.append(finding3)
        word4 = re.compile(r'puppo\s')
        finding4 = word4.findall(df6A.text_section[text])
        label_list.append(finding4)
    except:
        empty_list.append('')
```

```
In [77]: print('doggo:', label_list.count(['doggo ']))
print('floofer:', label_list.count(['floofer ']))
print('pupper:', label_list.count(['pupper ']))
print('puppo:', label_list.count(['puppo ']))
```

```
doggo: 40
floofer: 0
pupper: 75
puppo: 5
```

```
In [78]: # Let's compare the figures above to the number of times they appears
         in the labels column.
         df6A['dog_labels'].value_counts()
```

```
Out[78]:
pupper      1688
doggo       203
doggo       63
puppo       22
doggo and pupper    9
floofer      7
doggo and puppo    1
doggo and floofer  1
Name: dog_labels, dtype: int64
```

Conclusion: These terms appear more often in the dog labels column than they do in the actual text.

## Data Analysis, Part 4: Dog Breeds and Image Predictions

The neural network used to predict the dog breeds often predicted Chihuahua, Labrador retriever, and golden retriever. Let's compare the accuracy for each of those three breeds on the first two attempts.

```
In [79]: print(df6A.query("p1 == 'Chihuahua'")['p1_dog'].value_counts())
         print(df6A.query("p1 == 'Labrador_retriever'")['p1_dog'].value_counts(
         ))
         print(df6A.query("p1 == 'golden_retriever'")['p1_dog'].value_counts())

True      79
Name: p1_dog, dtype: int64
True      95
Name: p1_dog, dtype: int64
True     139
Name: p1_dog, dtype: int64
```

```
In [80]: print(df6A.query("p2 == 'Chihuahua'")['p2_dog'].value_counts())
print(df6A.query("p2 == 'Labrador_retriever'")['p2_dog'].value_counts(
))
print(df6A.query("p2 == 'golden_retriever'")['p2_dog'].value_counts())

True      43
Name: p2_dog, dtype: int64
True      96
Name: p2_dog, dtype: int64
True      82
Name: p2_dog, dtype: int64
```

## Data Wrangling Write-up

Data cleaning steps 7 and 8 are imperfect in that they do not take into account some special cases. For example, one particular tweet shows a rating of 20/16 in the 'text' column. Perhaps ratings like this could be converted to have a denominator of 10 but that would require code that considers a wide variety of cases. Another alternative might be to simply eliminate certain rows where *both* the numerator and the denominator are outside of a normal range.

Cleaning item number 9 came to mind only after I realized that the 3 DataFrames could not be joined or merged unless there was consistency in the 'tweet\_id' column. It is important to note that the extract method used is good but it did not capture all tweet\_ids since there are rows in which the 'expanded\_urls' column did not have the tweet\_id in it. This means that certain data was excluded from our master\_dataset but it was a relatively limited number of rows that were lost.

For Cleaning item number 10 I ran into certain difficulties in extracting the dog's name from the text column in cases where it was not present in the 'name' column. Also, I should mention that there were three or four cases where the dog's name appears as a noun such as 'not' or 'officially'. These were isolated cases so they did not have an impact on Analysis item number 2, which dealt with dog names.

For the data tidiness section I had intended to use the "melt" method but the presence of multiple labels for the same dog in about 15 cases made it difficult to do so. The methods I used required more code but provided a clearer solution and kept the data complete.

The last item to note is that there were a number of cases where I found it difficult to use vectorization so I opted for a for-loop instead. I understand that this dataset was relatively small so this was not a problem. With a larger dataset then vectorization would be preferable.

## Other Conclusions and Observations

There were a number of unique cases within this dataset that were not always easy to account for unless 1) substantial amounts of code were written to account for those unique cases or 2) many manual updates were made to account for specific cases. This did not, however, have a major impact on the conclusions and observations that were made in the analysis section at the end of this report. Using 1994 unique tweets was sufficient to give a good idea of some of the averages and trends within the data.

There are some specific items that would require a larger dataset in order provide better credibility. For example, the data is adequate for some conclusions in Data Analysis item 4 but a dataset would likely allow for a larger number of predictions for each breed of dog and therefore lend better credibility to this analysis. I found it a little hard to believe that the model accurately predicted the dog breed 100% of the time for Chihuahua, Labrador retriever and golden retriever. Also, there seemed to be some contradiction; if the dog was accurately predicted to be one breed then how could it be also accurately deemed to be another? Either the dogs are often a mix of various breeds or this is an area for further investigation and explanation.